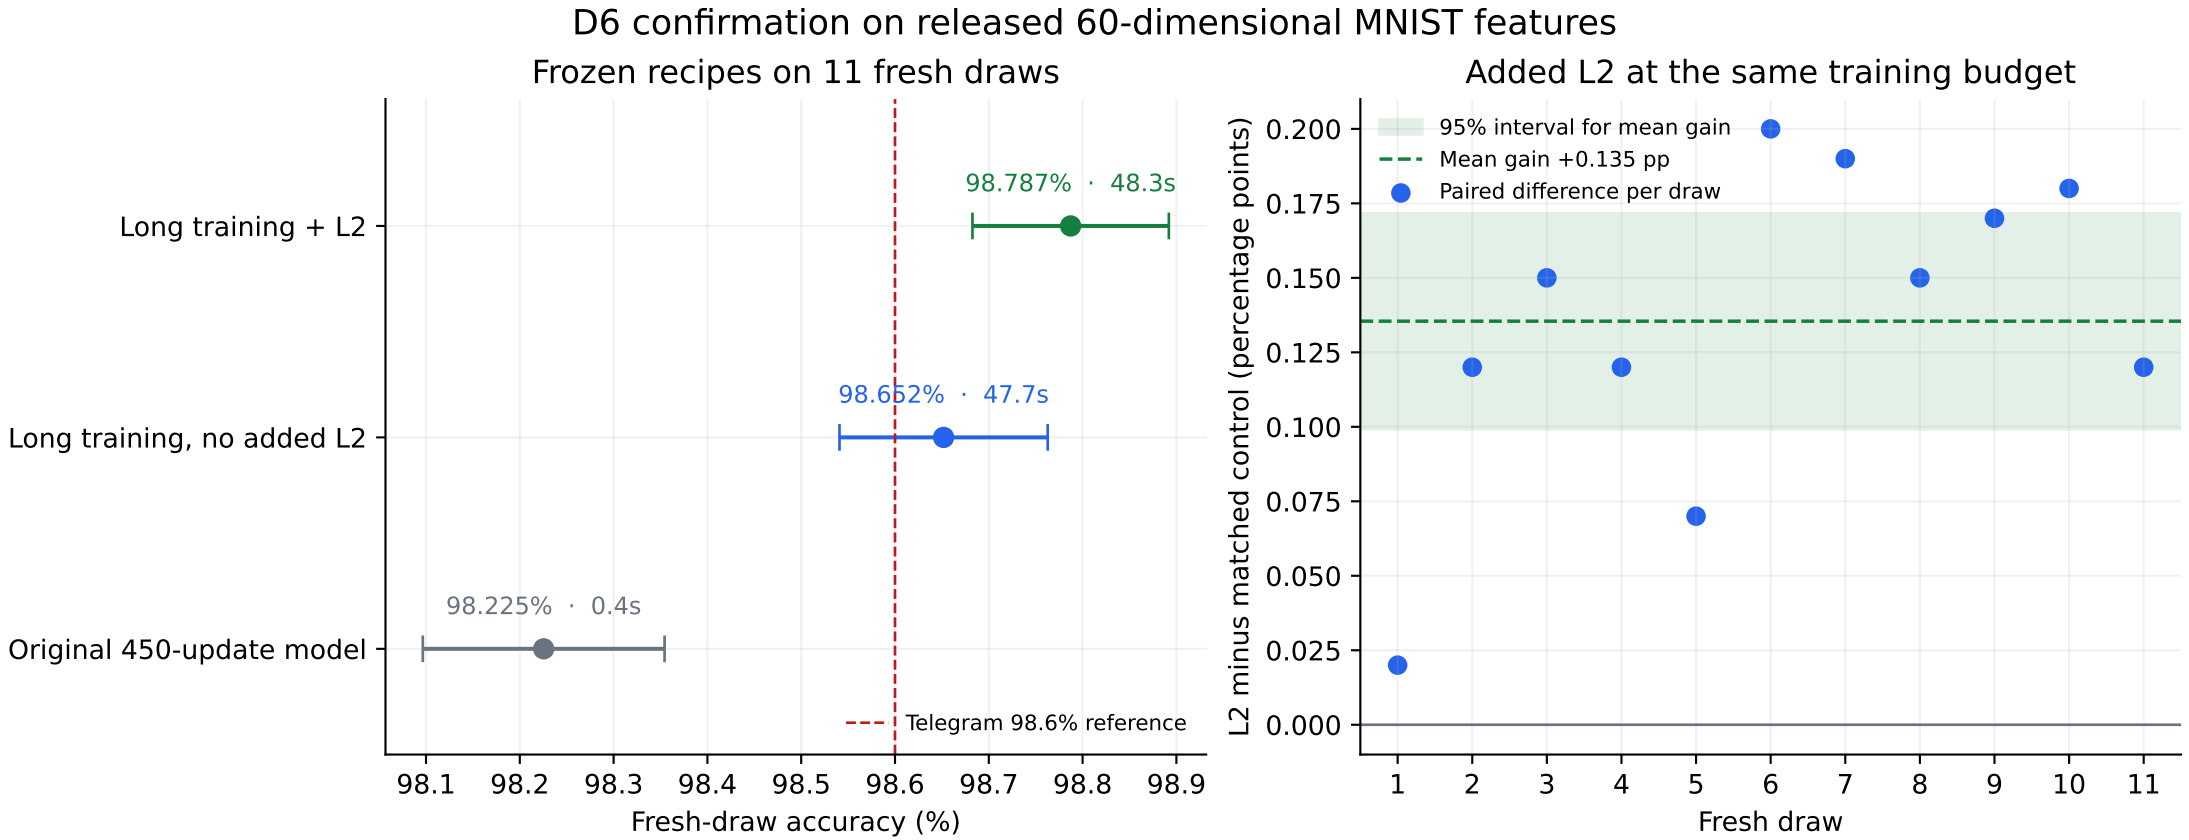
D6 confirmation on released 60-dimensional MNIST features
Frozen recipes on 11 fresh draws
Long training + L2
98.787% · 48.3s
Long training, no added L2
98.652% · 47.7s
Original 450-update model
98.225% · 0.4s
Telegram 98.6% reference
98.1
98.2
98.3
98.4
98.5
98.6
98.7
98.8
98.9
Fresh-draw accuracy (%)
Added L2 at the same training budget
95% interval for mean gain
Mean gain +0.135 pp
Paired difference per draw
L2 minus matched control (percentage points)
0.200
0.175
0.150
0.125
0.100
0.075
0.050
0.025
0.000
1
2
3
4
5
6
7
8
9
10
11
Fresh draw
Intervals describe draw-level means conditional on the fixed pool. Research timing excludes warmup and evidence serialization.